

DS002R - HW 7 - Web scraping

due Tuesday, November 12, 2024

Bernice Abanda

Q1 SongKick

Scrape the list of top 200 artists from https://www.songkick.com/leaderboards/popular_artists. (SongKick is a website designed to track artists and their live tour dates.)

Don't forget to change `#| eval: false` to `#| eval: true` so that your code will run (after you've filled in the blanks).

a. Read the html page

```
library(rvest)
#| eval: false
page <- read_html("/Users/berniceabanda/git/songkick_page.html")
```

b. Scrape the artists

```
artists <- page |>
  html_elements(".name strong") |>
  html_text()
```

c. Scrape the number of fans

```
library(stringr)
library(rvest)
library(httr)

fans <- page |>
  html_elements(".leaderboard-list-item .count") |>
  html_text() |>
  str_extract("\\d+") |>
  str_remove_all(",") |>
  as.numeric()
```

d. Scrape the number of concerts

```
concerts <- page |>
  html_elements(".leaderboard-list-item .count + .count") |>
  html_text() |>
  str_extract("\\d+") |>
  as.numeric()
```

e. Put it all in a data frame

```
songkick_top_200 <- tibble(
  artist = artists,
  fan = fans,
  concert = concerts
)
```

f. Add the rank associated with each artist. Print the first few row of the tibble.

```
songkick_top_200 <- songkick_top_200 |>
  mutate(rank = 0:nrow(songkick_top_200)) |>
  relocate(rank)
```

Q2 WNBA

Consider the data on WNBA standings for the last 15 years provided at <https://www.wnba.com/standings>.

a. Go to the WNBA standings website and scrape the standings table. (Note that the `html_table()` function returns a list of tables, even though there is only one table. To get the first table (one out of one), you'll need `html_table() |> pluck(1)` whose output should be a tibble.) Print the first few row of the tibble.

```
library(rvest)
library(tibble)
library(purrr)
library(ggplot2)
wnba_page <- read_html("https://www.wnba.com/standings")
wnba_standings <- wnba_page |>
  html_table() |>
  pluck(1)
head(wnba_standings)
```

```
# A tibble: 6 x 10
  TEAM      W      L  PCT GB  CONF HOME ROAD STREAK `L-10`
  <chr>    <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <chr>
1 1NYLClose    32     8 0.8  --   16-4  16-4  16-4  L1     7-3
2 2MINClose    30    10 0.75  2    14-6  16-4  14-6  L1     8-2
3 3CONClose    28    12 0.7   4    14-6  14-6  14-6  W1     6-4
4 4LVAClose    27    13 0.675 5    12-8  13-7  14-6  W5     9-1
5 5SEAClose    25    15 0.625 7    13-7  14-6  11-9  W1     6-4
6 6INDClose    20    20 0.5   12    11-9  12-8  8-12  L1     6-4
```

```
print(wnba_standings)
```

```
# A tibble: 12 x 10
  TEAM      W      L  PCT GB  CONF HOME ROAD STREAK `L-10`
  <chr>    <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <chr>
1 1NYLClose    32     8 0.8  --   16-4  16-4  16-4  L1     7-3
2 2MINClose    30    10 0.75  2    14-6  16-4  14-6  L1     8-2
3 3CONClose    28    12 0.7   4    14-6  14-6  14-6  W1     6-4
4 4LVAClose    27    13 0.675 5    12-8  13-7  14-6  W5     9-1
5 5SEAClose    25    15 0.625 7    13-7  14-6  11-9  W1     6-4
6 6INDClose    20    20 0.5   12    11-9  12-8  8-12  L1     6-4
7 7PHOClose    19    21 0.475 13   10-10 10-10 9-11  L1     3-7
8 8ATLClose    15    25 0.375 17    7-13  8-12  7-13  W3     5-5
9 9WASCircle    14    26 0.35  18    7-13  5-15  9-11  W1     6-4
10 10CHICircle  13    27 0.325 19    5-15  6-14  7-13  L5     2-8
11 11DALCircle   9    31 0.225 23    6-14  7-13  2-18  L9     1-9
12 12LASCircle   8    32 0.2   24    5-15  5-15  3-17  W1     2-8
```

- b. Clean up the `TEAM` variable so that it has two columns: first is the final `rank`, second is the three digit `team` code for each team. Print the first few row of the tibble.

```
wnba_standings <- wnba_standings |>
  separate(TEAM, into = c("rank", "team"), sep = " ", extra = "merge")
head(wnba_standings)
```

```
# A tibble: 6 x 11
  rank  team      W      L  PCT GB  CONF HOME ROAD STREAK `L-10`
  <chr> <chr> <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <chr>
1 1NYLClose <NA>    32     8 0.8  --   16-4  16-4  16-4  L1     7-3
2 2MINClose <NA>    30    10 0.75  2    14-6  16-4  14-6  L1     8-2
3 3CONClose <NA>    28    12 0.7   4    14-6  14-6  14-6  W1     6-4
```

4	4LVAClose	<NA>	27	13	0.675	5	12-8	13-7	14-6	W5	9-1
5	5SEAClose	<NA>	25	15	0.625	7	13-7	14-6	11-9	W1	6-4
6	6INDClose	<NA>	20	20	0.5	12	11-9	12-8	8-12	L1	6-4

c. Do the same thing for the 2014 season. Print the first few row of the tibble.

```
# Filter for 2014 season data

standings_2014 <- read_html("https://www.wnba.com/standings/?season=2014") |>
  html_table() |>
  pluck(1)

#Clean up the variables
standings_2014 <- standings_2014 |>
  mutate(
    rank = str_extract(Team, "\\d+"),
    team = str_remove(Team, "\\d+\\s"),
    Team = NULL
  )
# Print the first few rows of the tibble
print(standings_2014)
```

```
# A tibble: 12 x 11
      W     L   PCT GB   CONF HOME ROAD STREAK `L-10` rank team
  <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr>
1    19    15 0.559 --   11-11 13-4  6-11  L1    4-6    1    1ATLClose
2    29     5 0.853 --   19-3   16-1  13-4  W2    8-2    1    1PHOCClose
3    25     9 0.735 --   15-7   15-2  10-7  W1    7-3    2    2MINCClose
4    16    18 0.471 --   12-10  7-10  9-8   L1    5-5    2    2INDCClose
5    16    18 0.471 --    9-13   8-9   8-9   W3    4-6    3    3LVAClose
6    16    18 0.471 --   11-11   8-9   8-9   L1    5-5    3    3WASCClose
7    16    18 0.471 --    9-13   7-10  9-8   L1    6-4    4    4LASClose
8    15    19 0.441 --   14-8    9-8   6-11  L2    5-5    4    4CHICClose
9    15    19 0.441 --   10-12  10-7   5-12  W2    5-5    5    5NYLCCircle
10   12    22 0.353 --    7-15   8-9   4-13  L2    3-7    5    5SEACCircle
11   12    22 0.353 --    7-15   8-9   4-13  L3    4-6    6    6DALCCircle
12   13    21 0.382 --    8-14   9-8   4-13  W1    3-7    6    6CONCCircle
```

d. Write a function whose only argument is year and scrapes the standings for that year. The year argument should be a column in the output data frame. (That is, the function result is a dataframe with columns: **rank**, **team**, **W**, **L**, ... **L-10**, **year**. Remove the **TEAM** column after you've used it to create **rank** and **team** variables.) Print the first few row of the tibble.

```

library(rvest)
library(dplyr)
library(tidyr)
library(purrr)

# function to scrape the standings for a given year
scrape_standings <- function(year) {
  standings <- read_html(paste0("https://www.wnba.com/standings/?season=", year)) |>
    html_table() |>
    pluck(1)

  standings <- standings |>
    mutate(
      rank = str_extract(TEAM, "^\\d+"),
      team = str_remove(TEAM, "^\\d+\\s"),
      TEAM = NULL,
      year = year
    )

  return(standings)
}

```

- e. Use `map()` to scrape all the team data between 2010 and 2024. Print the first few rows of the tibble.

```

library(rvest)
library(dplyr)
library(tidyr)
library(purrr)

years <- 2010:2024
standings_all <- map(years, scrape_standings) |>
  list_rbind()

# Print the first few rows of the tibble
print(standings_all)

```

```

# A tibble: 180 x 12
      W      L    PCT GB   CONF HOME ROAD STREAK `L-10` rank team      year
  <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <int>
1    22    12 0.647 --   13-9  13-4  9-8   W6     8-2     1   1WASclose 2010
2    28     6 0.824 --   20-2  17-0 11-6   W3     6-4     1   1SEAClose 2010

```

```

3      22      12 0.647 --      14-8  13-4  9-8   W1      9-1    2      2NYLClose 2010
4      15      19 0.441 --      13-9   9-8  6-11  L2      4-6    2      2PHOClose 2010
5      14      20 0.412 --      11-11  8-9  6-11  W2      5-5    3      3LVAClose 2010
6      21      13 0.618 --      13-9   12-5  9-8   L3      5-5    3      3INDClose 2010
7      13      21 0.382 --      10-12  8-9  5-12  L1      5-5    4      4LASClose 2010
8      19      15 0.559 --      10-12  10-7  9-8   L2      4-6    4      4ATLClose 2010
9      17      17 0.5    --      9-13   12-5  5-12  L1      4-6    5      5CONCirc~ 2010
10     13      21 0.382 --      8-14   7-10  6-11  W1      5-5    5      5MINCirc~ 2010
# i 170 more rows

```

- f. Using `ggplot()`, create a line plot with `year` on the x-axis and `PCT` on the y-axis. Color each line according to the `team`. Do you see any trends of the teams over time?

```

library(ggplot2)

ggplot(standings_all, aes(x = year, y = PCT, color = team)) +
  geom_line() +
  labs(title = "WNBA Team Performance Over Time", x = "Year", y = "Win Percentage") +
  theme_minimal()

```

3A Team Performance Over Time

Q3 Best places

Go to <https://www.bestplaces.net> and search for Claremont, California. The website is used for comparing cities to determine where you might want to work or live.

- a. Using the SelectorGadget, extract the following pieces of information from the Claremont page:
 - property crime (just the number which is on a scale of 0 to 100)
 - minimum income required for a single person to live comfortably (as a number)
 - average monthly rent for a 2-bedroom apartment (as a number)
 - the “about” paragraph (the very first paragraph above “Location Details”)

Print the first few row of the tibble.

```
# Load necessary libraries
library(rvest)
library(tibble)
library(stringr)

# Read the html page
page <- read_html("https://www.bestplaces.net/city/california/claremont")

# Scrape the property crime rate
crime_rate <- page |>
  html_elements(".crime_rate") |>
  html_text() |>
  str_extract("\\d+") |>
  as.numeric()

# Scrape the minimum income required for a single person to live comfortably
min_income <- page |>
  html_elements(".min_income") |>
  html_text() |>
  str_extract("\\d+") |>
  as.numeric()

# Scrape the average monthly rent for a 2-bedroom apartment
rent_2br <- page |>
  html_elements(".rent_2br") |>
  html_text() |>
  str_extract("\\d+") |>
  as.numeric()
```

```
# Scrape the "about" paragraph
about <- page |>
  html_elements(".about") |>
  html_text()

# Put it all in a data frame
best_places <- tibble(
  state = "California",
  city = "Claremont",
  crime = crime_rate,
  min_income_single = min_income,
  rent_2br = rent_2br,
  about = about
)

# Print the first few rows of the tibble
print(best_places)
```

```
# A tibble: 0 x 6
# i 6 variables: state <chr>, city <chr>, crime <dbl>, min_income_single <dbl>,
#   rent_2br <dbl>, about <chr>
```

- b. Write a function called `scrape_bestplaces()` with arguments for city and state. When you run, for example `scrape_bestplaces("california", "claremont")`, the output should be a 1x6 tibble with columns for `state`, `city`, `crime`, `min_income_single`, `rent_2br`, and `about`.

```
scrape_bestplaces <- function(state, city) {
  page <- read_html(paste0("https://www.bestplaces.net/city/", state, "/", city))

  crime <- page |>
    html_element(".crime-index .value") |>
    html_text() |>
    as.numeric()

  min_income <- page |>
    html_element(".income-index .value") |>
    html_text() |>
    str_remove_all("[^0-9]") |>
    as.numeric()

  rent <- page |>
```



```

    html_element(".rent-index .value") |>
    html_text() |>
    str_remove_all("[^0-9]") |>
    as.numeric()

about <- page |>
  html_element(".about p") |>
  html_text()

tibble(
  state = state,
  city = city,
  crime = crime,
  min_income_single = min_income,
  rent_2br = rent,
  about = about
)
}

```

- c. Using `map2()`, create a 5 x 6 tibble by running `scrape_bestplaces()` 5 times with 5 cities you are interested in. Be sure you look at the URL at bestplaces.net for the various cities to make sure it works as you expect. Print the first few row of the tibble.

Running the `map2()` function will look something like this:

```

states <- c("california", "new-york", "texas", "florida", "oregon")
cities <- c("claremont", "new-york-city", "austin", "miami", "portland")

city_data <- map2(states, cities, scrape_bestplaces) |>
  list_rbind()

head(city_data)

```

```

# A tibble: 5 x 6
  state      city      crime min_income_single rent_2br about
  <chr>    <chr>    <dbl>         <dbl>      <dbl> <chr>
1 california claremont      NA              NA        NA <NA>
2 new-york   new-york-city  NA              NA        NA <NA>
3 texas      austin         NA              NA        NA <NA>
4 florida    miami          NA              NA        NA <NA>
5 oregon     portland       NA              NA        NA <NA>

```

Q4 Permission

Check to make sure we are allowed to scrape data from the three websites: songkick, wnba, and bestplaces.